

Xinyang Han

hanxinyang@berkeley.edu | Website | Scholar | GitHub | LinkedIn | 📍 Berkeley, CA

AI systems researcher building verifiable evaluations for long-horizon agents; co-lead of ALE and core contributor to MLS-Bench, with prior work in multimodal learning and 3D vision.

Education

University of California, Berkeley

PhD Student | Berkeley Artificial Intelligence Research (BAIR)

Berkeley, CA

Sep. 2024 – Present.

University of Chinese Academy of Sciences

Computer Science Major

Concurrent Enrollment Student in Computer Science, UC Berkeley

Beijing, China

2020–2024

Jan.–Sep. 2023

Selected Research

Agents' Last Exam (ALE)

Co-Lead

Co-leading a living benchmark that tests whether AI agents can reliably perform economically valuable work across real professional domains, not just short or synthetic tasks.

ALE converts expert-completed projects into long-horizon GUI+CLI workflows with verifiable outcomes; built with **300+ experts**, it spans **1,500+ tasks** across **55 sub-industries** and was used by OpenAI in the GPT-5.6 launch.

JobBench: Aligning Agent Work With Human Will

Co-author

Co-authored a benchmark of workflows workers most want delegated to AI, shifting workplace-agent evaluation from replacement value to human demand.

Covers **130 tasks** across **35 occupations**, grounded in ratings from **1,500+ workers**; **36 agent configurations** were evaluated with fact-anchored rubric chains, with the strongest reaching **45.9%**.

MLS-Bench: Can AI Systems Build Better AI?

Core Contributor

Core contributor to a benchmark testing whether agents can invent ML methods that generalize across controlled settings and scales.

Across **140 tasks** in **12 domains**, experiments show that tuning is easier than method invention and that more search, compute, or context alone does not remove the scientific-insight bottleneck.

SAP3D: Sparse-View 3D Reconstruction

Co-first Author

Co-developed a system that reconstructs 3D objects from sparse, unposed images by jointly estimating camera poses and adapting a view-conditioned diffusion model at test time.

Accepted as a **CVPR 2024 Highlight**.

U2Seg: Universal Unsupervised Image Segmentation

Co-first Author

Co-developed a unified framework for unsupervised instance, semantic, and panoptic segmentation.

Accepted to **CVPR 2024**.

AI for Cancer Prediction

PhD Research

Researching multimodal methods and evaluation for cancer risk prediction at BAIR, UC Berkeley.

Focus: clinically grounded, high-stakes decision support with physician collaborators.

Leadership & Service

Founder, AI Investment Summit 2025 @ UC Berkeley

Founded and led a student-run AI conference with **1,000+ attendees**; managed a **100+ person team** and assembled a 20+ speaker program spanning frontier AI labs and venture firms.

Academic service

Organizer, Workshop on Agent Behavior at COLM 2026; Communications Chair, CHIL 2026; Logistics Chair, CHIL 2025.

Reviewer

ICLR 2025, NeurIPS 2025, KDD 2025, and CHIL 2025.

Selected Publications

- Yiyou Sun, **Xinyang Han**, Weichen Zhang, Yuanbo Pang, Tianyu Wang, Yuhan Cao, Yixiao Huang, et al. *Agents' Last Exam*. arXiv, 2026. **Co-Lead**.
- Yuetai Li, Yichen Feng, Zhangchen Xu, Zixian Ma, ..., **Xinyang Han**, et al. *JobBench: Aligning Agent Work With Human Will*. arXiv, 2026.
- Bohan Lyu, Yucheng Yang, Siqiao Huang, Jiaru Zhang, Qixin Xu, Xinghan Li, **Xinyang Han**, et al. *MLS-Bench: A Holistic and Rigorous Assessment of AI Systems on Building Better AI*. **Core Contributor**. ICML 2026 AI for Math Workshop.
- Xinyang Han***, Zelin Gao*, Angjoo Kanazawa, Shubham Goel, Yossi Gandelsman. *The More You See in 2D, the More You Perceive in 3D*. **CVPR 2024 Highlight**.
- Dantong Niu*, Xudong Wang*, **Xinyang Han***, Long Lian, Roei Herzig, Trevor Darrell. *Unsupervised Universal Image Segmentation*. CVPR 2024.
- Yujie Yuan*, **Xinyang Han***, Yue He, Fanglue Zhang, Lin Gao. *MuNeRF: Robust Makeup Transfer in Neural Radiance Fields*. IEEE TVCG, 2025.
- Shu-Yu Chen, Yue-Ren Jiang, Hongbo Fu, **Xinyang Han**, Zitao Liu, Rong Li, Lin Gao. *DeepFaceReshaping: Interactive Deep Face Reshaping via Landmark Manipulation*. Computational Visual Media, 2024 (CVM 2023).
- Tianyu Zhang, Weiqing Min, **Xinyang Han**, Shuqiang Jiang, Yong Rui. *A Review on Action Prediction in Videos*. Chinese Journal of Computers, 2023.

* Equal contribution.

Technical Skills

Evaluation systems	computer-use environments, long-horizon workflows, task and rubric design, artifact-based grading, model and harness experiments
Languages / frameworks	Python, C++, PyTorch, JAX, TensorFlow
Multimodal	representation learning, computer vision, sparse-view 3D reconstruction, universal segmentation

Invited Talks

- Agents' Last Exam: Evaluating AI Agents on Real Work*, Snorkel AI, UC Berkeley, Kimi AI, WebAgentLab, and Valkyrie AI, June–July 2026.
- 3D Foundation Models for Generation and Reconstruction*, ICT Turing Class, July 2024.