Xinyang Han

hanxinyang@berkeley.edu | Website | Scholar | GitHub | LinkedIn | Berkeley, CA

AI systems researcher building verifiable evaluations for long-horizon agents; core contributor to ALE and MLS-Bench, with prior work in multimodal learning and 3D vision.

Education

UCSF–UC Berkeley | Computational Precision Health — **Berkeley, CA**
Joint PhD Student | Yala Lab — *Sep. 2024 – Present.*
AI for Cancer Prediction | Berkeley Artificial Intelligence Research (BAIR)

Selected Research

Agents' Last Exam (ALE) — ***Core Contributor***

Building Berkeley RDI's living benchmark of long-horizon professional workflows, with artifact-level verification across shell, files, web, and GUI environments.

ALE now spans **1,500+ tasks** across **55 occupational subdomains**, developed with **300+ experts**; OpenAI used ALE in its GPT-5.6 launch evaluation of long-horizon professional work.

JobBench: Aligning Agent Work With Human Will — ***Co-author***

Co-authored a benchmark of workflows workers most want delegated to AI, shifting workplace-agent evaluation from replacement value to human demand.

Covers **130 tasks** across **35 occupations**, grounded in ratings from **1,500+ workers**; **36 agent configurations** were evaluated with fact-anchored rubric chains, with the strongest reaching **45.9%**.

MLS-Bench: Can AI Systems Build Better AI? — ***Core Contributor***

Core contributor to a benchmark testing whether agents can invent ML methods that generalize across controlled settings and scales.

Across **140 tasks** in **12 domains**, experiments show that tuning is easier than method invention and that more search, compute, or context alone does not remove the scientific-insight bottleneck.

SAP3D: Sparse-View 3D Reconstruction — ***Co-first Author***

Co-developed a system that reconstructs 3D objects from sparse, unposed images by jointly estimating camera poses and adapting a view-conditioned diffusion model at test time.

Accepted as a **CVPR 2024 Highlight**.

U2Seg: Universal Unsupervised Image Segmentation — ***Co-first Author***

Co-developed a unified framework for unsupervised instance, semantic, and panoptic segmentation.

Accepted to **CVPR 2024**.

AI for Cancer Prediction — ***PhD Research***

Researching multimodal methods and evaluation for cancer risk prediction in Yala Lab within the UCSF–UC Berkeley Computational Precision Health program.

Focus: clinically grounded, high-stakes decision support with physician collaborators.

Leadership & Service

Founder, AI Investment Summit 2025 @ UC Berkeley

Founded and led a student-run AI conference with **1,000+ attendees**; managed a **100+ person team** and assembled a 20+ speaker program spanning frontier AI labs and venture firms.

Academic service	Organizer, Workshop on Agent Behavior at COLM 2026; Communications Chair, CHIL 2026; Logistics Chair, CHIL 2025.
Reviewer	ICLR 2025, NeurIPS 2025, KDD 2025, and CHIL 2025.
Invited talk	*3D Foundation Models for Generation and Reconstruction*, ICT Turing Class, July 2024.

Selected Publications

Yiyou Sun, **Xinyang Han**, Weichen Zhang, Yuanbo Pang, Tianyu Wang, Yuhan Cao, Yixiao Huang, et al. *Agents' Last Exam*. arXiv, 2026. **Core Contributor.**

Yuetai Li, Yichen Feng, Zhangchen Xu, Zixian Ma, …, **Xinyang Han**, et al. *JobBench: Aligning Agent Work With Human Will*. arXiv, 2026.

Bohan Lyu, Yucheng Yang, Siqiao Huang, Jiaru Zhang, Qixin Xu, Xinghan Li, **Xinyang Han**, et al. *MLS-Bench: A Holistic and Rigorous Assessment of AI Systems on Building Better AI*. **Core Contributor.** ICML 2026 AI for Math Workshop.

Xinyang Han*, Zelin Gao*, Angjoo Kanazawa, Shubham Goel, Yossi Gandelsman. *The More You See in 2D, the More You Perceive in 3D*. **CVPR 2024 Highlight**.

Dantong Niu*, Xudong Wang*, **Xinyang Han***, Long Lian, Roei Herzig, Trevor Darrell. *Unsupervised Universal Image Segmentation*. CVPR 2024.

Yujie Yuan*, **Xinyang Han***, Yue He, Fanglue Zhang, Lin Gao. *MuNeRF: Robust Makeup Transfer in Neural Radiance Fields*. IEEE TVCG, 2025.

Shu-Yu Chen, Yue-Ren Jiang, Hongbo Fu, **Xinyang Han**, Zitao Liu, Rong Li, Lin Gao. *DeepFaceReshaping: Interactive Deep Face Reshaping via Landmark Manipulation*. Computational Visual Media, 2024 (CVM 2023).

Tianyu Zhang, Weiqing Min, **Xinyang Han**, Shuqiang Jiang, Yong Rui. *A Review on Action Prediction in Videos*. Chinese Journal of Computers, 2023.

** Equal contribution.*

Technical Skills

Evaluation systems	computer-use environments, long-horizon workflows, task and rubric design, artifact-based grading, model and harness experiments
Languages / frameworks	Python, C++, PyTorch, JAX, TensorFlow
Multimodal	representation learning, computer vision, sparse-view 3D reconstruction, universal segmentation